


Deep Learning in Python: Image Recognition for Anime Characters with Transfer Learning

1st PyCon in Indonesia - 2017

Iskandar Setiadi





Github

<https://github.com/freedomofkeima>

Website

<https://freedomofkeima.com/>

From Jakarta, Indonesia
Graduated from ITB - 2015
Speaker in PyCon JP - 2017



Iskandar Setiadi

HDE, Inc. (<https://hde.co.jp/en/>)


Software Engineer at Japan



Why Python?

- Easy to use
- Great community
- Swiss army knife: website development, data science, etc

Image Recognition in Daily Life




Self-driving car



Smart home

Image Recognition in Daily Life



Smart workplace



Face ID

Background



Problem

- Cropped images
- Edited images
- Unindexed images

The screenshot shows a search result for a cropped image of two girls in pink dresses. The top banner indicates the project is 56% funded. The main result is for "Illumination" (Pixiv ID: 61127204) by Member: tucana. Below the main result, there are additional details: Creator: tucana, Material: love live! school idol project, love live! sunshine!! love live! (series), and Characters: sakurauchi riko, nishikino maki. A message at the bottom says "Low similarity results have been hidden. Click here to display them...".

The screenshot shows a search result for an edited image of two girls in pink dresses. The top banner indicates the project is 56% funded. The main result is for "Shanimuni Go" (Shanimuni Go - v16 c91 [batoto] - (Manga)) by Member: 碧乃魅沙@雪燃オンリーE12. Below the main result, there is a note: "雪燃オンリー兄さんといつしょ【新刊】雪燃合同誌：週末監禁".

Problem

- Photos

← → C | 🔒 Secure | <https://www.iqdb.org>

Multi-service image search

Search results

[Main page](#)

Searched 14,281,000 images in 1.128 seconds.

Your image	No relevant matches
 download.jpeg 750×1000	Could not find your image on any of the selected services. Maybe try this search on: SauceNao Google Images TinEye Below you can find the images that are the most similar to yours.


Possible match  Anime-Pictures 1644×1165 [Safe] 47% similarity	Possible match  Zerochan 592×902 [Safe] 47% similarity	Possible match  Zerochan 1644×1165 [Ero] 45% similarity
--	--	---

ILSVRC


Largest Computer Vision Competition


Starting from 2015, deep learning has better top-5 error score compared to human (1000 categories)!

Case #2: ILSVRC 2010-2015



Revolution of Depth






Tutorial for ML Beginner: MNIST & TensorFlow

55000 Training data

5000 Validation data

10000 Test data




~

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.6	.8	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	.1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	.1	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.5	.1	.4	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.4	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.4	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.7	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.1	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.9	.1	.1	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.3	.1	.1	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

URL: https://www.tensorflow.org/get_started/mnist/beginners

Tutorial for ML Beginner: MNIST & TensorFlow



URL: https://www.tensorflow.org/get_started/mnist/beginners



TensorFlow Installation

```
$ pip3 install --upgrade tensorflow
```

or

```
$ pip3 install --upgrade tensorflow-gpu
```

URL: <https://www.tensorflow.org/install/>



MNIST Model: TensorFlow + Python

```
x = tf.placeholder(tf.float32, [None, 784]) # Placeholder  
W = tf.Variable(tf.zeros([784, 10])) # Weight (W)  
b = tf.Variable(tf.zeros([10])) # Bias (b)  
  
# Tensor Flow it!  
# We can run it in CPU and GPU (let TensorFlow handle it)  
y = tf.nn.softmax(tf.matmul(x, W) + b)
```

MNIST Result & Comparison


Multilayer Neural Network with Logistic Regression Acc. : ~ 91%
Speed (1000 iter, 0.01 learning rate): < 1 minute

Convolutional Neural Network (Deep Learning) Acc.: ~ 99%
Speed (20000 iter, 0.0001 learning rate):
~2700 seconds (without GPU), ~360 seconds (with GPU)

Deep Learning

Increasing number of iterations will get stagnated at certain point.


More layers! But it is slow :(



Deep Learning: Convolution

1 x1	1 x0	1 x1	0	0
0 x0	1 x1	1 x0	1	0
0 x1	0 x0	1 x1	1	1
0	0	1	1	0
0	1	1	0	0


Image



Convolved Feature


0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter




Visualization of a curve detector filter


Deep Learning: Convolution



Original image




Visualization of the filter on the image



Visualizations of filters


Face Detection: Introduction



Face Detection (Human Face)

Adapted from <https://github.com/shantnu/FaceDetect>:

Face Detection (Human Face)




Face Detection (Same Model - Anime Face)



2D is Better not equal to 3D face!

Facial features are different!

e.g.: 2D has no nose



Face Detection: Train New Model!

Adapted from https://github.com/nagadomi/lbpcascade_animeface:

Face Detection (Anime Face)



Face Detection (Same Model - Human Face)






Face Recognition

Face Detection → “Accomplished”

Full-layered Deep Learning → Requires a huge dataset, weeks to train

Google Inception-v3: 1.2 million training data, 1000 classes, 1 week to train

What does the layers learn?



- Orange square: Convolution
- Blue square: AvgPool
- Green square: MaxPool
- Red square: Concat
- Purple square: Dropout
- Pink square: Fully connected
- Dark red square: Softmax



Transfer Learning


From certain Top-5 characters indexing website:

- 35000 registered characters
- Top 1000 characters: 70+ images
- Top 2000 characters: 40+ images

Dataset size is small! Google Inception-v3 uses > 1000 images per category.

With transfer learning, we don't need to retrain low-level features extraction model.

URL: https://www.tensorflow.org/tutorials/image_retraining



Transfer Learning: Retrained Layers

Dropout: Dropping out units to prevent overfitting

Fully Connected: Extracting global features, every node in the layer is connected to the preceding layer

Softmax: Squashing final layer to make a prediction, which sums up to 1. For example, if we have 2 classes and class A has the value of 0.95, then class B will have the value of 0.05.

Transfer Learning: Retrain Final Layer

Build the retrainer:

```
$ bazel build tensorflow/examples/image_retraining:retrain
```


Execute the retrainer:

```
$ bazel-bin/tensorflow/examples/image_retraining/retrain --image_dir ~/images
```


Hyperparameters: learning rate, number of iterations, distortions factor, ...

Types of headache


MIGRAINE




HYPERTENSION



STRESS



Working with classifiers





MoeFlow: Web App

```
@app.listener('before_server_start')
async def initialize(app, loop):
    moeflow_path = os.environ.get('MOEFLOW_MODEL_PATH')
    label_path = os.path.join(os.sep, moeflow_path, "output_labels_2.txt")
    model_path = os.path.join(os.sep, moeflow_path, "output_graph_2.pb")
    app.label_lines = [
        line.strip() for line in tf.gfile.GFile(label_path)
    ]
    graph = tf.Graph()
    graph_def = tf.GraphDef()
    with tf.gfile.FastGFile(model_path, 'rb') as f:
        graph_def.ParseFromString(f.read())
    with graph.as_default():
        tf.import_graph_def(graph_def, name='')
    app.graph = graph
    logging.info("MoeFlow model is now initialized!")
```

URL: <https://github.com/freedomofkeima/MoeFlow>



MoeFlow: Specification

- Build with Sanic (Flask-like Python 3.5+ web server)
- While training model requires huge GPU resources (g2.2xlarge), using retrained model can be hosted in server with small resources (t2.micro)

What it does:

- Run face detection with OpenCV
- Resize image to a fixed proportion
- Run classification with TensorFlow

```
def classify_resized_face(file_name, label_lines, graph):
    results = []
    logging.info('Processing classification')
    with tf.Session(graph=graph) as sess:
        # Feed the image data as input to the graph and get first prediction
        softmax_tensor = sess.graph.get_tensor_by_name('final_result:0')
        input_operation = sess.graph.get_operation_by_name("Mul")
        t = read_tensor_from_image_file(file_name)
        predictions = sess.run(
            softmax_tensor,
            {input_operation.outputs[0]: t}
        )
        # Sort to show labels of first prediction in order of confidence
        top_k = predictions[0].argsort()[-3:][::-1]

        for node_id in top_k:
            human_string = label_lines[node_id]
            score = predictions[0][node_id]
            results.append((human_string, score))
    return results
```




```
2017-12-06 01:51:53 - (network) [INFO] [127.0.0.1:49154]: GET http://127.0.0.1:888/static/images/83a97094d7634333b5c9d7e8e0901da4.jpg 200 119582
INFO:network:
2017-12-06 01:51:53 - (network) [INFO] [127.0.0.1:49156]: GET http://127.0.0.1:888/static/images/e96410fcf8114fbb88a4f33ede40205e.jpg 200 5902
INFO:network:
2017-12-06 05:48:33 - (network) [INFO] [127.0.0.1:59920]: GET http://127.0.0.1:888/ 200 1180
INFO:network:
INFO:root:Height: 1416, Width: 1280
INFO:root:Input file is created at /tmp/tmpgmg8odxy2.jpg
INFO:root:Processing classification
INFO:root:[('tedeza rize', 0.8941825), ('yui', 0.015284972), ('takanashi rikka', 0.014930859)]
2017-12-06 05:49:38 - (network) [INFO] [127.0.0.1:59970]: POST http://127.0.0.1:888/ 200 1565
INFO:network:
2017-12-06 05:49:39 - (network) [INFO] [127.0.0.1:59972]: GET http://127.0.0.1:888/static/images/fda2e4333b194f30a845ed8d320bf449.jpg 200 294605
INFO:network:
2017-12-06 05:49:39 - (network) [INFO] [127.0.0.1:59974]: GET http://127.0.0.1:888/static/images/d9a530d76c17404fbdbde7a2e5bf9c55.jpg 200 5807
INFO:network:
```

Test Results (Number of Class)

With 100 class and 60 images per class, it achieves 70.1% top-1 accuracy.

When the number of class is relatively small (~35), it can achieve 80%+ top-1 accuracy.



URL: https://github.com/freedomofkeima/MoeFlow/blob/master/100_class_traning_note.md



Test Results (Dataset size)

100 class experiment:

- 30 images per class: **60.3%** accuracy
- 60 images per class: **70.1%** accuracy

All tests are done with images which are not in training / validation set.

URL: https://github.com/freedomofkeima/MoeFlow/blob/master/100_class_traning_note.md



Output:



Prediction: tedeza rize, kasumigaoka utaha, hyoudou michiru



Prediction: kirima sharo, hoto cocoa, alice cartelet



Prediction: hoto cocoa, kirima sharo, yagami kou

```
ec2-user[ip-172-31-48-241:~/MoeFlow] - □ ×
2017-12-02 10:38:45 - (network) [INFO][127.0.0.1:35584]: POST http://127.0.0.1:8888/
200 1577
INFO:network:
2017-12-02 10:38:45 - (network) [INFO][127.0.0.1:35586]: GET http://127.0.0.1:8888/static/images/bdb9c730b1074a4c8506fc7e0d733c7d.jpg 200 160714
INFO:network:
2017-12-02 10:38:45 - (network) [INFO][127.0.0.1:35588]: GET http://127.0.0.1:8888/static/images/f16e3eb03197475c865a0237fe0c7db7.jpg 200 6142
INFO:network:
INFO:root:Height: 1574, Width: 2048
INFO:root:Input file is created at /tmp/tmppjelu9cy8.jpg
INFO:root:Processing classification
INFO:root:[('tedeza rize', 0.39948133), ('kasumigaoka utaha', 0.12605736), ('hyoudou michiru', 0.068339251)]
INFO:root:Processing classification
INFO:root:[('kirima sharo', 0.97146475), ('hoto cocoa', 0.018899804), ('alice cartelet', 0.0034658003)]
INFO:root:Processing classification
INFO:root:[('hoto cocoa', 0.94451982), ('kirima sharo', 0.028552018), ('yagami kou', 0.0080380896)]
2017-12-02 10:39:08 - (network) [INFO][127.0.0.1:35590]: POST http://127.0.0.1:8888/
200 2013
INFO:network:
2017-12-02 10:39:08 - (network) [INFO][127.0.0.1:35592]: GET http://127.0.0.1:8888/
```

Note: This operation will be very slow (around 15 seconds) if there are a lot of characters in a single image!

No file chosen

Input:



```
ec2-user@ip-172-31-48-241:~/Moeflow
2017-12-02 11:25:42 - (network) [INFO] [127.0.0.1:38124]: GET http://127.0.0.1:888
8/static/images/3775c305912b482693c7f32daaf52c1.jpg 200 6643
INFO:network:
2017-12-02 11:25:42 - (network) [INFO] [127.0.0.1:38126]: GET http://127.0.0.1:888
8/static/images/4cca249295f440leb45fc0e34da809c2.jpg 200 7175
INFO:network:
2017-12-02 11:25:42 - (network) [INFO] [127.0.0.1:38128]: GET http://127.0.0.1:888
8/static/images/8e6db243c7aa458c98061de61c0496c6.jpg 200 6126
INFO:network:
INFO:root:Height: 800, Width: 534
INFO:root:Input file is created at /tmp/tmp135gbz41.jpg
INFO:root:Processing classification
INFO:root:[('kousaka kirino', 0.4233555), ('sakura nene', 0.35790154), ('dekomori sanae', 0.08103735)]
2017-12-02 11:26:55 - (network) [INFO] [127.0.0.1:38178]: POST http://127.0.0.1:888/ 200 1575
INFO:network:
2017-12-02 11:26:56 - (network) [INFO] [127.0.0.1:38180]: GET http://127.0.0.1:888
8/static/images/dc40cb23b9534196a3226f9f843e8500.jpg 200 80962
INFO:network:
2017-12-02 11:26:56 - (network) [INFO] [127.0.0.1:38182]: GET http://127.0.0.1:888
8/static/images/a2dc5e325342434f898495ec5b8cf64f.jpg 200 5796
INFO:network:
```

Output:



Prediction: kousaka kirino, sakura nene, dekomori sanae

Input:



Output:



Prediction: suzukaze aoba, kafuu chino, yamada elf



Prediction: yagami kou, ayase eli, tomo mami

Rotation / Axis problem


```
ec2-user@ip-172-31-48-241:~/MoeFlow
2017-12-02 11:00:43 - (network) [INFO] [127.0.0.1:36820]: GET http://127.0.0.1:8888/static/images/d6bfe2062fa844d88e7095cff241eb40.jpg 200 5372
INFO:network:
INFO:root:Height: 860, Width: 1280
INFO:root:Input file is created at /tmp/tmpish217do.jpg
INFO:root:Processing classification
INFO:root:[('suzukaze aoba', 0.90099531), ('kafuu chino', 0.030359417), ('yamada elf', 0.013034027)]
INFO:root:Processing classification
INFO:root:[('yagami kou', 0.89891928), ('ayase eli', 0.043274842), ('tomo mami', 0.026960159)]
2017-12-02 11:00:59 - (network) [INFO] [127.0.0.1:36822]: POST http://127.0.0.1:8888/ 200 1782
INFO:network:
2017-12-02 11:00:59 - (network) [INFO] [127.0.0.1:36824]: GET http://127.0.0.1:8888/static/images/02dfd53aeafc42d4b0189bd376f0f05b.jpg 200 264179
INFO:network:
2017-12-02 11:00:59 - (network) [INFO] [127.0.0.1:36826]: GET http://127.0.0.1:8888/static/images/7c65786b9db74bdbaa83a97f908ff9823.jpg 200 6871
INFO:network:
2017-12-02 11:01:00 - (network) [INFO] [127.0.0.1:36828]: GET http://127.0.0.1:8888/static/images/7a968786c4a24f01a8dfd998e4ba363e.jpg 200 6488
INFO:network:
```

Problems (Example)

“Never-ending” Development

- Image noise
- Rotation / axis
- Face expressions (closed eyes, etc)
- Characters with “multiple” forms
- Brightness & Contrast

Fooling Neural Network



Original image
Output Label: Teapot



Noisy image (10% impulse noise)
Output Label: Biology

Image Recognition as a Service

If you need image recognition features for production-ready environment and you don't have any specific requirements to build your model from ground:

- Amazon Rekognition
- Computer Vision API in Cognitive Service (Azure)

Amazon Rekognition

Deep learning-based image recognition service
Search, verify, and organize millions of images



Object and Scene
Detection



Facial
Analysis



Face
Comparison



Facial
Recognition

Amazon Rekognition

Metrics

Demos

Object and scene detection

Image moderation

Facial analysis

Celebrity recognition

Face comparison

Text in image

Video Demos

Video analysis

Additional Resources

Getting started guide

Download SDKs

Developer resources

Pricing

FAQ

Forum

Face comparison

Compare faces to see how closely they match based on a similarity percentage.



InvalidParameterException (400)


Request has Invalid Parameters (Image must contain detectable faces)

Reference face



Choose a sample image

Comparison faces



Choose a sample image

Done with the demo?

[Learn more](#)

▶ Results

▶ Request

▶ Response



My Github Projects

[freedomofkeima/MoeFlow](#): Repository for anime characters recognition website (Alpha)

[freedomofkeima/transfer-learning-anime](#): Transfer Learning for Anime Characters Recognition

[freedomofkeima/opencv-playground](#): Compare 2D and 3D OpenCV Cascade Classifier

Presentation Slide

<https://freedomofkeima.com/pyconid2017.pdf>

Curated List

<https://github.com/kjw0612/awesome-deep-vision>

<http://www.themtank.org/a-year-in-computer-vision>

HDE, Inc. at Shibuya, Tokyo

- Global Internship Program (<https://www.hde.co.jp/en/gip/>)
- 15% international people
- 6 people from Indonesia





Thank you!